

Healthcare Analytics Pipeline for MediCare Hospital

Submitted By:

Arpit Sharma — Roll No: G25AI1012

Ankita Dwivedi — Roll No: G25AI1011

Astosh Ranjan — Roll No: G25AI1013

Bharat Sharma — Roll No: G25AI1014

Damini Khatri — Roll No: G25AI1015

Course: Post Graduate Diploma in Data Engineering (PGDDE)

Institute: IIT Jodhpur

Plagiarism Declaration

We declare that this project report, implementation, configurations, scripts, workflows, dashboards, database designs, and deployment procedures are the original work of our team carried out during academic practice.

All necessary references, documentation sources, and publicly available resources used during the implementation have been properly acknowledged. Any form of plagiarism may lead to disciplinary action as per academic policies.

PART 1 — OBJECTIVE

The objective of this project is to:

- Build an end-to-end Healthcare Analytics Pipeline.
- Automate healthcare data ingestion, processing, and reporting.
- Store and manage healthcare data using cloud technologies.
- Perform large-scale healthcare analytics using PySpark.
- Design a healthcare data warehouse using Star Schema.
- Develop interactive dashboards for healthcare insights.
- Implement machine learning for predictive healthcare analytics.

This project demonstrates the practical implementation of Data Engineering, Cloud Computing, Big Data Processing, Data Warehousing, Analytics, and Visualization technologies in a healthcare environment.

PART 2 — PROBLEM STATEMENT

Healthcare organizations generate large volumes of patient, treatment, and operational data every day.

Traditional approaches:

- Require manual data processing and reporting.
- Are time-consuming and error-prone.
- Lack centralized data management.
- Struggle with large-scale data processing.
- Provide limited analytical capabilities.

This system addresses these challenges by:

- Automating data ingestion and processing.
- Centralizing healthcare data storage.
- Enabling scalable analytics using cloud infrastructure.
- Providing automated workflow execution.
- Delivering interactive dashboards and business insights.

PART 3 — DATASET DESCRIPTION

The project uses a healthcare dataset containing patient records and treatment information.

Dataset Characteristics:

- Total Records: 5000
- Data Format: CSV
- Domain: Healthcare Analytics
- Storage Platform: AWS S3

Dataset Attributes:

- Patient ID
- Patient Name
- Age
- Gender
- Disease
- City
- Treatment Cost
- Admission Status

- Admission Date

The dataset contains anonymized healthcare information and was used for ETL processing, analytics, data warehousing, dashboard development, and machine learning implementation.

PART 4 — TECHNOLOGY STACK

The following technologies were used during project implementation:

Component	Technology Used
Programming Language	Python
Data Processing	Pandas
Workflow Automation	Apache Airflow
Big Data Processing	PySpark
Cloud Storage	AWS S3
Database	PostgreSQL
Cloud Database	AWS RDS PostgreSQL
Data Warehouse	Star Schema
Dashboard	Streamlit
Machine Learning	Scikit-Learn
Version Control	Git & GitHub
Containerization	Docker

The selected technology stack provides scalability, automation, cloud integration, and analytical capabilities required for modern healthcare data engineering solutions.

PART 5 — IMPLEMENTATION

The Healthcare Analytics Pipeline was implemented using a modular architecture consisting of cloud storage, ETL processing, workflow automation, database integration, big data processing, analytics, and visualization components.


Step 1 — AWS S3 Data Storage

Healthcare records were stored in AWS S3 cloud storage before processing.

Benefits:

- Centralized data storage

- Cloud-based accessibility
- Scalability for future data growth
- Integration with ETL workflows



Step 2 — ETL Pipeline Implementation

An ETL (Extract, Transform, Load) pipeline was developed using Python and Pandas.

Activities Performed:

- Dataset extraction
- Data cleaning
- Missing value handling
- Data validation
- Data transformation
- Loading into PostgreSQL

The ETL pipeline automated healthcare data processing and reduced manual effort.

```
(base) PS C:\Users\arpit\OneDrive\Desktop\healthcare-analytics-pipeline> cd scripts
(base) PS C:\Users\arpit\OneDrive\Desktop\healthcare-analytics-pipeline\scripts> python etl_pipeline.py
Healthcare Dataset Loaded Successfully
patient_id    name    age    gender    disease    admission_date    discharge_date    hospital_bill    city    blood_pressure    sugar_level    heart_rate
0    P001    Rahul Sharma    45    Male    Diabetes    2026-05-01    2026-05-05    25000    Delhi    140/90    180    88
1    P002    Priya Verma    32    Female    Hypertension    2026-05-02    2026-05-06    18000    Mumbai    150/95    110    92
2    P003    Aman Gupta    60    Male    Heart Disease    2026-05-03    2026-05-10    75000    Bangalore    160/100    140    105
3    P004    Neha Singh    28    Female    Asthma    2026-05-04    2026-05-07    12000    Chennai    120/80    95    80
4    P005    Rohan Das    50    Male    Diabetes    2026-05-05    2026-05-09    30000    Kolkata    145/92    210    90

Data Cleaning Completed

PostgreSQL Connection Successful

Data Loaded into PostgreSQL Successfully
(base) PS C:\Users\arpit\OneDrive\Desktop\healthcare-analytics-pipeline\scripts>
```



Step 3 — Apache Airflow Automation

Apache Airflow was implemented for workflow orchestration and automation.

Features:

- DAG-based workflow execution
- Automated scheduling
- Monitoring and logging
- Pipeline management

The Airflow DAG successfully automated ETL execution and workflow monitoring.



Step 4 — PySpark Big Data Processing

Apache Spark was integrated into the pipeline using PySpark to process large healthcare datasets efficiently.

Activities Performed:

- Loading healthcare dataset into Spark
- Schema validation
- Record counting
- Disease distribution analysis
- Treatment cost analysis

Results Obtained:

- Successfully processed 5000 healthcare records
- Generated disease-wise analytics
- Demonstrated scalable distributed processing

```
(base) PS C:\Users\arpit\OneDrive\Desktop\healthcare-analytics-pipeline\scripts> python spark_pipeline.py
Healthcare Big Dataset Loaded Using Spark

First 10 Records
-----+-----+-----+-----+-----+-----+-----+-----+
|patient_id|patient_name|age|gender|disease|city|treatment_cost|admission_status|admission_date|
-----+-----+-----+-----+-----+-----+-----+-----+
|1|Jorge Smith MD|79|Male|HEART DISEASE|Hyderabad|139410|Admitted|11/2/2025|
|2|Mary Ryan|28|Male|HYPERTENSION|Chennai|162031|Discharged|6/22/2025|
|3|Jacob Rodriguez|34|Female|ASTHMA|Chennai|132177|Discharged|4/9/2026|
|4|Patrick Santana|2|Female|DIABETES|Bangalore|133709|Discharged|8/3/2025|
|5|Taylor Le|37|Male|COVID|Pune|105963|Admitted|9/4/2025|
|6|Pamela Pope|31|Female|COVID|Hyderabad|98968|Discharged|3/27/2026|
|7|Antonio Alvarado|46|Female|HYPERTENSION|Bangalore|51897|Admitted|10/13/2025|
|8|Brenda Morgan|31|Male|HEART DISEASE|Pune|37838|Discharged|12/5/2025|
|9|Angela Gonzalez|18|Male|CANCER|Bangalore|98497|Discharged|7/17/2025|
|10|Anna Nunez|49|Female|CANCER|Mumbai|176029|Admitted|7/3/2025|
-----+-----+-----+-----+-----+-----+-----+-----+

only showing top 10 rows

Dataset Schema
root
 |-- patient_id: integer (nullable = true)
 |-- patient_name: string (nullable = true)
 |-- age: integer (nullable = true)
 |-- gender: string (nullable = true)
 |-- disease: string (nullable = true)
 |-- city: string (nullable = true)
 |-- treatment_cost: integer (nullable = true)
 |-- admission_status: string (nullable = true)
 |-- admission_date: string (nullable = true)

Total Patients: 5000

Disease Distribution
-----+-----+
|disease|count|
-----+-----+
|HYPERTENSION|827|
|CANCER|840|
|DIABETES|845|
|HEART DISEASE|805|
|ASTHMA|861|
|COVID|822|
-----+-----+

Average Treatment Cost
-----+-----+
|avg(treatment_cost)|
-----+-----+
|101866.402|
-----+-----+

PySpark Processing Completed Successfully
(base) PS C:\Users\arpit\OneDrive\Desktop\healthcare-analytics-pipeline\scripts> SUCCESS: The process with PID 18820 (child process of PID 29328) has been terminated.
SUCCESS: The process with PID 29328 (child process of PID 30760) has been terminated.
SUCCESS: The process with PID 30760 (child process of PID 4396) has been terminated.
```


Step 5 — AWS RDS PostgreSQL Integration

AWS RDS PostgreSQL was used as the cloud database for storing processed healthcare records.

Benefits:

- Managed database service
- Cloud-based accessibility
- Secure data storage
- High availability and reliability

The ETL pipeline successfully loaded healthcare records into AWS RDS PostgreSQL.



Step 6 — Data Warehouse Design

A Star Schema Data Warehouse was implemented to support analytical reporting and dashboard development.

Warehouse Components:

Dimension Tables

- dim_patient
- dim_disease
- dim_date

Fact Table

- fact_patient_visits

Benefits:

- Faster analytical queries
- Better reporting performance
- Simplified data analysis
- Improved dashboard responsiveness

```
(base) PS C:\Users\arpit\OneDrive\Desktop\healthcare-analytics-pipeline\scripts> python check_warehouse.py

Data Warehouse Tables

dim_date
dim_disease
dim_patient
fact_patient_visits
patients

Warehouse Check Completed Successfully
(base) PS C:\Users\arpit\OneDrive\Desktop\healthcare-analytics-pipeline\scripts>
```


Step 7 — Interactive Dashboard Development

A Streamlit dashboard was developed for healthcare analytics and reporting.

Dashboard Features:

- Total Patients KPI
- Average Treatment Cost
- Disease Distribution Analysis
- Patient Data Visualization
- Interactive Reporting

The dashboard provides healthcare insights through a user-friendly interface.



PART 6 — RESULTS

The Healthcare Analytics Pipeline was successfully implemented and validated.

Major Achievements:

- Processed 5000 healthcare records successfully
- Automated ETL workflow using Python
- Integrated AWS S3 cloud storage
- Implemented Apache Airflow automation
- Performed distributed processing using PySpark
- Integrated AWS RDS PostgreSQL
- Designed a Star Schema Data Warehouse
- Developed a Streamlit analytics dashboard
- Implemented machine learning analytics

Performance Summary

Component	Status
AWS S3 Integration	Successful
ETL Pipeline	Successful
Airflow Automation	Successful
PySpark Processing	Successful
AWS RDS Integration	Successful
Data Warehouse	Successful
Dashboard Development	Successful
Machine Learning	Successful

```
(base) PS C:\Users\arpit\OneDrive\Desktop\healthcare-analytics-pipeline\scripts> python ml_prediction.py
Dataset Size: 5000 Records
Mean Absolute Error: 53185.24
```

Sample Predictions:

Actual Cost: ₹57876	Predicted Cost: ₹158560.59
Actual Cost: ₹65388	Predicted Cost: ₹128636.86
Actual Cost: ₹173984	Predicted Cost: ₹128576.27
Actual Cost: ₹94836	Predicted Cost: ₹95609.71
Actual Cost: ₹23624	Predicted Cost: ₹95532.52
Actual Cost: ₹194891	Predicted Cost: ₹121762.7
Actual Cost: ₹97918	Predicted Cost: ₹125949.52
Actual Cost: ₹12579	Predicted Cost: ₹121660.76
Actual Cost: ₹178794	Predicted Cost: ₹100741.79
Actual Cost: ₹99612	Predicted Cost: ₹123201.61

Model Training Completed Successfully

```
(base) PS C:\Users\arpit\OneDrive\Desktop\healthcare-analytics-pipeline\scripts> █
```

```
=====
EDA REPORT
=====
(base) PS C:\Users\arpit\OneDrive\Desktop\healthcare-analytics-pipeline\scripts> python eda_analysis.py

Total Records:
5000

Columns:
['patient_id', 'patient_name', 'age', 'gender', 'disease', 'city', 'treatment_cost', 'admission_status', 'admission_date']

Missing Values:
patient_id      0
patient_name    0
age             0
gender          0
disease         0
city            0
treatment_cost  0
admission_status 0
admission_date  0
dtype: int64

Disease Distribution:
disease
ASTHMA      861
DIABETES    845
CANCER      840
HYPERTENSION 827
COVID       822
HEART DISEASE 805
Name: count, dtype: int64

Gender Distribution:
gender
Male    2534
Female  2466
Name: count, dtype: int64

City Distribution:
city
Chennai    863
Mumbai     839
Hyderabad  834
Delhi      830
Bangalore  828
Pune       806
Name: count, dtype: int64

Admission Status Distribution:
admission_status
Admitted    2511
Discharged  2489
Name: count, dtype: int64

Average Treatment Cost:
101866.4

Maximum Treatment Cost:
199989

Minimum Treatment Cost:
5004


EDA Completed Successfully
```

PART 7 — COMPARISON AND ANALYSIS

Feature	Traditional Approach	Proposed System
Data Processing	Manual	Automated
Storage	Local Files	AWS Cloud
Workflow Management	Manual	Airflow
Big Data Support	Limited	PySpark
Reporting	Static	Interactive Dashboard
Scalability	Low	High
Analytics	Basic	Advanced
Machine Learning	Not Available	Implemented

The proposed system significantly improves scalability, automation, reporting capabilities, and analytical performance compared to traditional healthcare data processing approaches.

PART 8 — ARCHITECTURE DIAGRAM



Healthcare Analytics Pipeline Architecture

This architecture represents the complete flow of data from ingestion and processing to analytics, machine learning, and dashboard visualization.

PART 9 — DATA GOVERNANCE

Healthcare data requires proper governance, validation, and security controls to ensure data quality and reliability.

Data Governance Measures:

- Validation of incoming healthcare records
- Missing value detection and handling
- Duplicate record identification
- Data consistency checks
- Structured database design
- Secure cloud storage using AWS services
- Controlled access to healthcare information

Benefits:

- Improved data quality
 - Reliable analytics
 - Better reporting accuracy
 - Enhanced data integrity
-

PART 10 — CONCLUSION

This project successfully demonstrates the implementation of a complete Healthcare Analytics Pipeline for MediCare Hospital.

The system integrates AWS S3, Apache Airflow, PySpark, PostgreSQL, AWS RDS, Data Warehousing, Machine Learning, and Streamlit into a single analytics platform.

Major Achievements:

- Automated healthcare data processing
- Cloud-based storage and database integration
- Big data processing using PySpark
- Interactive dashboard development
- Machine learning implementation

- Data warehouse design for analytics

The project demonstrates practical applications of Data Engineering concepts and provides a scalable solution for healthcare analytics.

PART 11 — FUTURE SCOPE

The Healthcare Analytics Pipeline can be further enhanced using advanced technologies and real-time processing capabilities.

Future Enhancements:

- Real-time data streaming using Apache Kafka
- Advanced machine learning models
- Integration with hospital management systems
- Automated alert generation
- Cloud-native deployment
- Mobile dashboard integration
- Predictive healthcare analytics
- Enterprise-scale implementation

These improvements can further increase the scalability and analytical capabilities of the platform.

PART 12 — SOURCE CODE STRUCTURE

Project Directory Structure:

healthcare-analytics-pipeline/

├── dashboards/

| └─ app.py

├── data/

| └─ patients.csv

| └─ healthcare_big_data.csv

├── docker_airflow/

| └─ dags/

| └─ healthcare_etl_dag.py

├── docs/

└─ screenshots/

```
└─ scripts/
| └─ etl_pipeline.py
| └─ spark_pipeline.py
| └─ data_warehouse.py
| └─ ml_prediction.py
| └─ eda_analysis.py
└─ sql/
| └─ healthcare_queries.sql
└─ README.md
└─ requirements.txt
└─ .gitignore
```

PART 13 — ISSUES FACED & SOLUTIONS

Issue	Solution
Airflow Webserver Not Starting	Reconfigured Docker Compose Setup
AWS RDS Connectivity Issues	Corrected Database Configuration
ETL File Path Errors	Updated Relative File Paths
PySpark Dataset Loading Issues	Modified Dataset References
Dashboard Rendering Issues	Updated Streamlit Components
Data Warehouse Loading Errors	Corrected Table Structure and Schema

The above issues were resolved through testing, debugging, and configuration improvements.

PART 14 — DEMO & LINKS

GitHub Repository

The complete project source code, documentation, ETL scripts, Airflow DAGs, PySpark jobs, SQL queries, data warehouse implementation, and dashboard code are available on GitHub.

GitHub Repository Link:

<https://github.com/arpit-systems/healthcare-analytics-pipeline>

Live Dashboard

The Streamlit dashboard provides healthcare analytics, KPI monitoring, disease analysis, and interactive reporting.

Dashboard Link:

<https://medicare-healthcare-dashboard.streamlit.app>

Project Demonstration

Project Workflow:

Healthcare Dataset → AWS S3 → ETL Pipeline → Apache Airflow → PySpark → AWS RDS PostgreSQL → Data Warehouse → Machine Learning → Streamlit Dashboard

The project successfully automates healthcare data processing, analytics, and reporting through an end-to-end data engineering pipeline.

PART 15 — TEAM CONTRIBUTION

This project was completed through the collaborative efforts of all team members. Responsibilities included data collection, data cleaning, ETL development, workflow automation using Apache Airflow, cloud integration, PySpark processing, data warehousing, machine learning, dashboard development, testing, documentation, presentation preparation, and project review.

All team members actively contributed to the successful completion of the Healthcare Analytics Pipeline for MediCare Hospital project.

FINAL NOTE

This report presents the complete implementation of a Healthcare Analytics Pipeline developed using modern Data Engineering technologies. The project demonstrates cloud integration, workflow automation, big data processing, analytics, machine learning, and interactive visualization within a healthcare domain.

The implementation successfully fulfills the objectives of the capstone project and provides a scalable foundation for future healthcare analytics applications.